

**FACULTY OF ENGINEERING AND TECHNOLOGY (Co-Ed)**  
**DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING**



**LAB MANUAL**

# **Neural Networks and Deep Learning Laboratory**

**(22CSL78)**

**2025-2026**

<p align="center"><b>Neural Networks and Deep Learning Laboratory</b></p> <p align="center"><b>[As per Choice Based Credit System(CBCS)scheme]</b></p> <p align="center"><b>(Effective from the academic year 2025-2026)</b></p> <p align="center"><b>SEMESTER – VII</b></p>			
<b>Course Code</b>	<b>22CSL78</b>	<b>CIE Marks</b>	<b>50</b>
<b>Number of Lecture Hours/Week</b>	<b>02</b>	<b>SEE Marks</b>	<b>50</b>
<b>Total Number of Lecture Hours</b>	<b>30</b>	<b>Exam Hours</b>	<b>3</b>
<b>CREDIT– 01</b>			
<b>Course Objectives :</b> This course will enable students to:			
<ul style="list-style-type: none"> <li>• Implement and train various types of neural network models, starting with a basic multilayer perceptron (MLP) for classification tasks.</li> <li>• Apply deep learning to computer vision.</li> <li>• Use deep learning for natural language processing (NLP).</li> <li>• Explore advanced deep learning architectures.</li> <li>• Develop real-world deep learning applications.</li> </ul>			
<b>Programs</b>			
1)Create neural network class and initialize those weights and biases.			
2)Implement all activation functions in Neural Network.			
3)Implement Loss function for Neural Network.			
4)Implement Forward propagation and Backward propagation.			
5)Program to train and test a neural network.			
6)Train and test the Convolution neural network using any data set , preprocess it.			
7)Implement the Convolution neural network for image classification.			
8)Train and test the recurrent neural network using any data set , preprocess it.			
9)Implement Facial recognition using neural network.			
10) Implement object detection using neural network.			

**COURSE OUTCOMES:** At the end of the course the student will be able to:

<b>CO1</b>	Make use of deep learning techniques to implement fundamental concepts of Neural Networks
<b>CO2</b>	Apply convolutional neural networks (CNNs) for computer vision tasks.
<b>CO3</b>	Develop and apply recurrent neural networks (RNNs) for sequential data.
<b>CO4</b>	Apply deep learning techniques for facial recognition & object detection
<b>CO5</b>	Prepare a well organized laboratory report detailing experimental procedure,reports

	PO1	PO2	PO3	PO4	PO5	PO6	PO7	PO8	PO9	PO 10	PO 11	PSO 1	PSO 2	PSO 3
<b>CO1</b>	3	3	2	2	2	-	-	-	2	2	-			
<b>CO2</b>	3	3	3	2	2	-	-	1	2	-	-			
<b>CO3</b>	3	3	2	2	2	-	-	-	-	-	-			
<b>CO4</b>	3	3	3	2	3	-	-	-	2	3	2			
<b>Avg</b>	<b>3.0</b>	<b>3</b>	<b>2.5</b>	<b>2</b>	<b>2.2</b>	<b>-</b>	<b>-</b>	<b>1</b>	<b>1.5</b>	<b>1.2</b>	<b>1</b>			

## Prog 1: Create neural network class and initialize those weights and biases.

```
import numpy as np

class NeuralNetwork:
    def __init__(self, layers):
        layers (list): Number of neurons in each layer.
            Example: [3, 5, 2] means:
                - 3 neurons in input layer
                - 5 neurons in hidden layer
                - 2 neurons in output layer

        self.layers = layers
        self.weights = []
        self.biases = []
        self.initialize_parameters()

    def initialize_parameters(self):
        """Initialize weights and biases with random values."""
        np.random.seed(42) # for reproducibility

        for i in range(1, len(self.layers)):
            # He initialization for weights
            weight_matrix = np.random.randn(self.layers[i], self.layers[i - 1]) * np.sqrt(2. /
self.layers[i - 1])
            bias_vector = np.zeros((self.layers[i], 1))

            self.weights.append(weight_matrix)
            self.biases.append(bias_vector)

    def summary(self):
```

```
print("Neural Network Architecture:")
for i in range(len(self.layers) - 1):
    print(f"Layer {i} -> Layer {i+1}: "
          f"Weights shape {self.weights[i].shape}, "
          f"Biases shape {self.biases[i].shape}")
if __name__ == "__main__":
    nn = NeuralNetwork([3, 5, 2])
    nn.summary()
```

### **Output :**

Neural Network Architecture:

Layer 0 -> Layer 1: Weights shape (5, 3), Biases shape (5, 1)

Layer 1 -> Layer 2: Weights shape (2, 5), Biases shape (2, 1)

## Prog 2: Implement all activation functions in Neural Network.

```
import numpy as np
```

```
class ActivationFunctions:
```

```
    @staticmethod
```

```
    def sigmoid(x):
```

```
        return 1 / (1 + np.exp(-x))
```

```
    @staticmethod
```

```
    def sigmoid_derivative(x):
```

```
        s = ActivationFunctions.sigmoid(x)
```

```
        return s * (1 - s)
```

```
    @staticmethod
```

```
    def tanh(x):
```

```
        return np.tanh(x)
```

```
    @staticmethod
```

```
    def tanh_derivative(x):
```

```
        return 1 - np.tanh(x)**2
```

```
    @staticmethod
```

```
    def relu(x):
```

```
        return np.maximum(0, x)
```

```
    @staticmethod
```

```
    def relu_derivative(x):
```

```
        return np.where(x > 0, 1, 0)
```

```
    @staticmethod
```

```
    def leaky_relu(x, alpha=0.01):
```

```
        return np.where(x > 0, x, alpha * x)
```

```
@staticmethod
```

```
def leaky_relu_derivative(x, alpha=0.01):
```

```
    dx = np.ones_like(x)
```

```
    dx[x < 0] = alpha
```

```
    return dx
```

```
@staticmethod
```

```
def elu(x, alpha=1.0):
```

```
    return np.where(x > 0, x, alpha * (np.exp(x) - 1))
```

```
@staticmethod
```

```
def elu_derivative(x, alpha=1.0):
```

```
    return np.where(x > 0, 1, ActivationFunctions.elu(x, alpha) + alpha)
```

```
@staticmethod
```

```
def softmax(x):
```

```
    exps = np.exp(x - np.max(x, axis=0, keepdims=True)) # stability fix
```

```
    return exps / np.sum(exps, axis=0, keepdims=True)
```

```
@staticmethod
```

```
def softmax_derivative(x):
```

```
    s = ActivationFunctions.softmax(x)
```

```
    return np.diagflat(s) - np.dot(s, s.T)
```

```
@staticmethod
```

```
def linear(x):
```

```
    return x
```

```
@staticmethod
```

```
def linear_derivative(x):
```

```
    return np.ones_like(x)
```

```
if __name__ == "__main__":
```

```
x = np.array([-2.0, -1.0, 0.0, 1.0, 2.0])

print("Input:", x)
print("ReLU:", ActivationFunctions.relu(x))
print("Sigmoid:", ActivationFunctions.sigmoid(x))
print("Tanh:", ActivationFunctions.tanh(x))
print("Leaky ReLU:", ActivationFunctions.leaky_relu(x))
print("ELU:", ActivationFunctions.elu(x))
print("Softmax:", ActivationFunctions.softmax(x))
```

### **Output:**

```
Input: [-2. -1.  0.  1.  2.]
ReLU: [0.  0.  0.  1.  2.]
Sigmoid: [0.11920292 0.26894142 0.5      0.73105858 0.88079708]
Tanh: [-0.96402758 -0.76159416  0.      0.76159416  0.96402758]
Leaky ReLU: [-0.02 -0.01  0.   1.   2. ]
ELU: [-0.86466472 -0.63212056  0.     1.     2.     ]
Softmax: [0.01165623 0.03168492 0.08612854 0.23412166 0.63640865]
```



### Prog 3: Implement Loss function for Neural Network.

```
import numpy as np
```

```
class LossFunctions:
```

```
    @staticmethod
```

```
    def mse(y_true, y_pred):
```

```
        """Mean Squared Error loss."""
```

```
        return np.mean((y_true - y_pred) ** 2)
```

```
    @staticmethod
```

```
    def mse_derivative(y_true, y_pred):
```

```
        return 2 * (y_pred - y_true) / y_true.size
```

```
    @staticmethod
```

```
    def mae(y_true, y_pred):
```

```
        return np.mean(np.abs(y_true - y_pred))
```

```
    @staticmethod
```

```
    def mae_derivative(y_true, y_pred):
```

```
        return np.sign(y_pred - y_true) / y_true.size
```

```
    @staticmethod
```

```
    def binary_cross_entropy(y_true, y_pred):
```

```
        eps = 1e-15 # avoid log(0)
```

```
        y_pred = np.clip(y_pred, eps, 1 - eps)
```

```
        return -np.mean(y_true * np.log(y_pred) + (1 - y_true) * np.log(1 - y_pred))
```

```
    @staticmethod
```

```

def binary_cross_entropy_derivative(y_true, y_pred):
    eps = 1e-15
    y_pred = np.clip(y_pred, eps, 1 - eps)
    return -(y_true / y_pred) + (1 - y_true) / (1 - y_pred)) / y_true.size

@staticmethod
def categorical_cross_entropy(y_true, y_pred):
    y_true: one-hot encoded true labels
    y_pred: predicted probabilities (softmax outputs)
    eps = 1e-15
    y_pred = np.clip(y_pred, eps, 1 - eps)
    return -np.mean(np.sum(y_true * np.log(y_pred), axis=0))

@staticmethod
def categorical_cross_entropy_derivative(y_true, y_pred):
    eps = 1e-15
    y_pred = np.clip(y_pred, eps, 1 - eps)
    return -(y_true / y_pred) / y_true.shape[0]

if __name__ == "__main__":
    y_true = np.array([[1, 0, 0], [0, 1, 0]]).T # one-hot encoded (3 classes, 2 samples)
    y_pred = np.array([[0.7, 0.2, 0.1], [0.1, 0.8, 0.1]]).T

    print("Mean Squared Error:", LossFunctions.mse(y_true, y_pred))

    print("Categorical Cross Entropy:", LossFunctions.categorical_cross_entropy(y_true,
y_pred))

```

### Output:

Mean Squared Error: 0.03333333333333334

Categorical Cross Entropy: 0.2899092476264711

## Prog 4: Implement Forward propagation and Backward propagation.

```
import numpy as np

# Activation function (Sigmoid) and its derivative
def sigmoid(x):
    return 1 / (1 + np.exp(-x))

def sigmoid_derivative(x):
    return x * (1 - x)

# Input dataset (4 samples, 2 features)
X = np.array([
    [0, 0],
    [0, 1],
    [1, 0],
    [1, 1]
])

# Output labels (XOR problem)
y = np.array([[0], [1], [1], [0]])

# Set random seed for reproducibility
np.random.seed(42)

# Initialize weights and biases
input_layer_size = 2
hidden_layer_size = 2
output_layer_size = 1
```

```

W1 = np.random.uniform(size=(input_layer_size, hidden_layer_size))
b1 = np.random.uniform(size=(1, hidden_layer_size))
W2 = np.random.uniform(size=(hidden_layer_size, output_layer_size))
b2 = np.random.uniform(size=(1, output_layer_size))

# Training parameters
learning_rate = 0.5
epochs = 10000

# Training loop
for epoch in range(epochs):
    # ---- Forward propagation ----
    hidden_input = np.dot(X, W1) + b1
    hidden_output = sigmoid(hidden_input)

    final_input = np.dot(hidden_output, W2) + b2
    predicted_output = sigmoid(final_input)

    # ---- Backward propagation ----
    error = y - predicted_output
    d_predicted_output = error * sigmoid_derivative(predicted_output)

    error_hidden = d_predicted_output.dot(W2.T)
    d_hidden_layer = error_hidden * sigmoid_derivative(hidden_output)

    # ---- Update weights and biases ----
    W2 += hidden_output.T.dot(d_predicted_output) * learning_rate
    b2 += np.sum(d_predicted_output, axis=0, keepdims=True) * learning_rate
    W1 += X.T.dot(d_hidden_layer) * learning_rate
    b1 += np.sum(d_hidden_layer, axis=0, keepdims=True) * learning_rate

```

```
# Print progress occasionally
if epoch % 2000 == 0:
    loss = np.mean(np.square(error))
    print(f"Epoch {epoch} - Loss: {loss:.6f}")

# Final output
print("\n Final predicted output after training:")
print(predicted_output)
```

### **Output :**

```
Epoch 0 - Loss: 0.324659
Epoch 2000 - Loss: 0.002546
Epoch 4000 - Loss: 0.000902
Epoch 6000 - Loss: 0.000537
Epoch 8000 - Loss: 0.000380
```

Final predicted output after training:

```
[[0.01890475]
 [0.98371361]
 [0.98369334]
 [0.01686123]]
```

### Prog 5 : Program to train and test a neural network.

```
import numpy as np

# Activation + derivative
def sigmoid(x): return 1 / (1 + np.exp(-x))
def sigmoid_deriv(a): return a * (1 - a)

# Neural Network
class NeuralNet:
    def __init__(self, lr=0.1):
        np.random.seed(42)
        self.W1 = np.random.randn(4, 2) # 2→4
        self.b1 = np.zeros((4, 1))
        self.W2 = np.random.randn(1, 4) # 4→1
        self.b2 = np.zeros((1, 1))
        self.lr = lr

    def forward(self, X):
        self.Z1 = np.dot(self.W1, X) + self.b1
        self.A1 = np.maximum(0, self.Z1) # ReLU
        self.Z2 = np.dot(self.W2, self.A1) + self.b2
        self.A2 = sigmoid(self.Z2) # Sigmoid
        return self.A2

    def backward(self, X, Y):
        m = Y.shape[1]
        dZ2 = (self.A2 - Y) * sigmoid_deriv(self.A2)
        dW2 = np.dot(dZ2, self.A1.T) / m
        dB2 = np.sum(dZ2, axis=1, keepdims=True) / m
        dA1 = np.dot(self.W2.T, dZ2)
```

```

dZ1 = dA1 * (self.A1 > 0)
dW1 = np.dot(dZ1, X.T) / m
dB1 = np.sum(dZ1, axis=1, keepdims=True) / m
# Update
self.W1 -= self.lr * dW1; self.b1 -= self.lr * dB1
self.W2 -= self.lr * dW2; self.b2 -= self.lr * dB2

def train(self, X, Y, epochs=2000):
    for i in range(epochs + 1):
        out = self.forward(X)
        loss = np.mean((Y - out)**2)
        self.backward(X, Y)
        if i % 200 == 0:
            print(f'Epoch {i:4d} | Loss: {loss:.5f}')

# --- Train & Test ---
X = np.array([[0,0,1,1],[0,1,0,1]]) # inputs
Y = np.array([[0,1,1,0]])          # outputs

nn = NeuralNet(lr=0.1)
nn.train(X, Y, epochs=4000)

preds = nn.forward(X)
print("\nPredictions:\n", preds.round(3))

```

### Output :

```

Epoch  0 | Loss: 0.26242
Epoch 200 | Loss: 0.24559
Epoch 400 | Loss: 0.22857
Epoch 600 | Loss: 0.19455

```

Epoch 800 | Loss: 0.15112  
Epoch 1000 | Loss: 0.12031  
Epoch 1200 | Loss: 0.09432  
Epoch 1400 | Loss: 0.07337  
Epoch 1600 | Loss: 0.05717  
Epoch 1800 | Loss: 0.04514  
Epoch 2000 | Loss: 0.03632  
Epoch 2200 | Loss: 0.02974  
Epoch 2400 | Loss: 0.02476  
Epoch 2600 | Loss: 0.02096  
Epoch 2800 | Loss: 0.01801  
Epoch 3000 | Loss: 0.01569  
Epoch 3200 | Loss: 0.01379  
Epoch 3400 | Loss: 0.01227  
Epoch 3600 | Loss: 0.01099  
Epoch 3800 | Loss: 0.00993  
Epoch 4000 | Loss: 0.00903

Predictions:

[[0.143 0.923 0.923 0.063]]



## **Prog 6: Train and test the Convolution neural network using any data set , preprocess it.**

```
# Import libraries
import tensorflow as tf
from tensorflow.keras import layers, models

# Load and preprocess data
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()

# Normalize pixel values (0–1) and reshape for CNN (samples, height, width, channels)
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0
x_train = x_train[..., None] # add channel dimension
x_test = x_test[..., None]

# Build CNN model
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax')
])

# Compile model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

```
# Train the model

print("Training CNN on MNIST dataset...")

model.fit(x_train, y_train, epochs=3, batch_size=64, validation_split=0.1, verbose=1)

# Evaluate on test set

test_loss, test_acc = model.evaluate(x_test, y_test, verbose=0)

print(f"\nTest accuracy: {test_acc * 100:.2f}%")
```

### Output:

Epoch 1/3

**844/844** ————— **10s** 11ms/step - accuracy: 0.8800  
- loss: 0.4101 - val\_accuracy: 0.9817 - val\_loss: 0.0681

Epoch 2/3

**844/844** ————— **9s** 10ms/step - accuracy: 0.9816 -  
loss: 0.0602 - val\_accuracy: 0.9867 - val\_loss: 0.0462

Epoch 3/3

**844/844** ————— **8s** 10ms/step - accuracy: 0.9864 -  
loss: 0.0425 - val\_accuracy: 0.9887 - val\_loss: 0.0397

Test accuracy: 98.90%

## **Prog 7: Implement the Convolution neural network for image classification.**

```
import tensorflow as tf

from tensorflow.keras import layers, models

from tensorflow.keras.datasets import cifar10

# Load and preprocess the dataset
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

# Normalize pixel values (0–1)
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0

# Convert labels to flat array
y_train = y_train.flatten()
y_test = y_test.flatten()

# Build CNN model
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax') # 10 classes
])

# Compile model
model.compile(optimizer='adam',
```

```

        loss='sparse_categorical_crossentropy',
        metrics=['accuracy'])

# Train model

print(" Training CNN on CIFAR-10 dataset...")
model.fit(x_train, y_train, epochs=5, batch_size=64, validation_split=0.1, verbose=1)

# Evaluate model

print("\n Evaluating on test data...")

test_loss, test_acc = model.evaluate(x_test, y_test, verbose=1)
print(f"\n Test accuracy: {test_acc * 100:.2f}%")

# Save model

# model.save("cnn_cifar10_model.h5")

```

### Output :

```

704/704 ————— 11s 14ms/step - accuracy: 0.3326
- loss: 1.8153 - val_accuracy: 0.5308 - val_loss: 1.2969

```

Epoch 2/5

```

704/704 ————— 9s 13ms/step - accuracy: 0.5502 -
loss: 1.2693 - val_accuracy: 0.5732 - val_loss: 1.1751

```

Epoch 3/5

```

704/704 ————— 9s 13ms/step - accuracy: 0.6129 -
loss: 1.0956 - val_accuracy: 0.6386 - val_loss: 1.0217

```

Epoch 4/5

```

704/704 ————— 9s 13ms/step - accuracy: 0.6570 -
loss: 0.9834 - val_accuracy: 0.6634 - val_loss: 0.9626

```

Epoch 5/5

```

704/704 ————— 9s 13ms/step - accuracy: 0.6837 -
loss: 0.9105 - val_accuracy: 0.6844 - val_loss: 0.8937

```

```

313/313 ————— 1s 3ms/step - accuracy: 0.6730 -
loss: 0.9409

```

Test accuracy: 67.00%

### **Prog 8: Train and test the recurrent neural network using any data set , preprocess it.**

```
import tensorflow as tf

from tensorflow.keras import layers, models
from tensorflow.keras.datasets import imdb
from tensorflow.keras.preprocessing import sequence

# Load Dataset (IMDB Movie Reviews)
max_words = 10000 # Only keep top 10k most common words
max_len = 200     # Cut reviews after 200 words
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=max_words)

# Preprocess Data (Pad sequences to same length)
x_train = sequence.pad_sequences(x_train, maxlen=max_len)
x_test = sequence.pad_sequences(x_test, maxlen=max_len)

# Build RNN model (LSTM)
model = models.Sequential([
    layers.Embedding(input_dim=max_words, output_dim=64, input_length=max_len),
    layers.LSTM(64),
    layers.Dense(1, activation='sigmoid')
])

# Compile model
model.compile(optimizer='adam',
              loss='binary_crossentropy',
              metrics=['accuracy'])

# Show model summary
model.summary()
```

```
# Train model

print("\n Training RNN (LSTM) on IMDB dataset...")

model.fit(x_train, y_train, epochs=3, batch_size=64, validation_split=0.2, verbose=1)

# Evaluate on test data

print("\n Evaluating on test data...")

test_loss, test_acc = model.evaluate(x_test, y_test, verbose=1)

print(f"\n Test Accuracy: {test_acc * 100:.2f}%")
```

**Output :**

**Model: "sequential\_2"**

Layer (type)	Output Shape	Param #
embedding (Embedding)	?	0 (unbuilt)
lstm (LSTM)	?	0 (unbuilt)
dense_4 (Dense)	?	0 (unbuilt)

**Total params:** 0 (0.00 B)

**Trainable params:** 0 (0.00 B)

**Non-trainable params: 0 (0.00 B)**

Training RNN (LSTM) on IMDB dataset...

Epoch 1/3

**313/313** ————— **25s** 74ms/step - accuracy: 0.6952  
- loss: 0.5481 - val\_accuracy: 0.8476 - val\_loss: 0.3683

Epoch 2/3

**313/313** ————— **24s** 75ms/step - accuracy: 0.8951  
- loss: 0.2680 - val\_accuracy: 0.8500 - val\_loss: 0.3421

Epoch 3/3

**313/313** ————— **24s** 75ms/step - accuracy: 0.9297  
- loss: 0.1940 - val\_accuracy: 0.8508 - val\_loss: 0.3468

Evaluating on test data...

**782/782** ————— **14s** 18ms/step - accuracy: 0.8423  
- loss: 0.3727

Test Accuracy: 84.36%

## **Prog 9: Implement Facial recognition using neural network.**

```
import tensorflow as tf

from tensorflow.keras import layers, models

from sklearn.datasets import fetch_lfw_people

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import LabelEncoder

import numpy as np


# Load the LFW (Labeled Faces in the Wild) dataset

print(" Loading dataset...")

faces = fetch_lfw_people(min_faces_per_person=70, color=True, resize=0.4)

X = faces.images / 255.0      # Normalize pixel values

y = faces.target

target_names = faces.target_names


# Split into training and testing sets

x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Reshape data for CNN

x_train = np.expand_dims(x_train, -1) if x_train.ndim == 3 else x_train

x_test = np.expand_dims(x_test, -1) if x_test.ndim == 3 else x_test


# Build CNN model

model = models.Sequential([

    layers.Conv2D(32, (3,3), activation='relu', input_shape=x_train.shape[1:]),

    layers.MaxPooling2D(2,2),

    layers.Conv2D(64, (3,3), activation='relu'),
```



```
layers.MaxPooling2D(2,2),
layers.Flatten(),
layers.Dense(128, activation='relu'),
layers.Dense(len(target_names), activation='softmax')
])
```

```
# Compile model
```

```
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

```
# Show summary
```

```
model.summary()
```

```
# Train model
```

```
print("\n Training CNN for facial recognition...")
```

```
model.fit(x_train, y_train, epochs=5, batch_size=32, validation_split=0.1, verbose=1)
```

```
# Evaluate on test data
```

```
print("\n Evaluating model...")
```

```
test_loss, test_acc = model.evaluate(x_test, y_test, verbose=1)
```

```
print(f"\nTest Accuracy: {test_acc * 100:.2f}%")
```

```
# Predict on one test sample
```

```
sample = np.expand_dims(x_test[0], axis=0)
```

```
pred = np.argmax(model.predict(sample), axis=1)[0]
```

```
print(f"\n Predicted: {target_names[pred]}")
```

```
print(f" Actual: {target_names[y_test[0]]}")
```

## Output :

Training CNN for facial recognition...

Epoch 1/5

28/28 [=====] - 6s 170ms/step - loss: 1.3391 - accuracy:  
0.5784 - val\_loss: 0.7633 - val\_accuracy: 0.8023

Epoch 5/5

28/28 [=====] - 5s 164ms/step - loss: 0.3412 - accuracy:  
0.9056 - val\_loss: 0.4232 - val\_accuracy: 0.8800

Evaluating model...

9/9 [=====] - 0s 22ms/step - loss: 0.4412 - accuracy:  
0.8760

Test Accuracy: 87.60%

## Prog 10: Implement object detection using neural network.

```
import tensorflow as tf

import cv2

import numpy as np


# Load a pre-trained object detection model from TensorFlow Hub

import tensorflow_hub as hub


print(" Loading pre-trained MobileNet SSD model...")

detector = hub.load("https://tfhub.dev/tensorflow/ssd_mobilenet_v2/2")


# Load labels for COCO dataset

labels_path = tf.keras.utils.get_file(

    'coco_labels.txt',

    'https://raw.githubusercontent.com/amikeli/coco-labels/master/coco-labels-paper.txt'

)

with open(labels_path, "r") as f:

    labels = [line.strip() for line in f.readlines()]


# Load and preprocess an image

image_path = tf.keras.utils.get_file(

    'dog.jpg',

    'https://storage.googleapis.com/mledu-

datasets/cats_and_dogs_filtered/train/dogs/dog.1.jpg'

)

image = cv2.imread(image_path)

image_rgb = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

input_tensor = tf.convert_to_tensor([image_rgb], dtype=tf.uint8)
```

```

# Run detection

print(" Running object detection...")

detections = detector(input_tensor)


# Extract detection results

boxes = detections['detection_boxes'][0].numpy()

classes = detections['detection_classes'][0].numpy().astype(int)

scores = detections['detection_scores'][0].numpy()


# Draw boxes for detections above threshold

for i in range(len(scores)):

    if scores[i] > 0.5:

        ymin, xmin, ymax, xmax = boxes[i]

        h, w, _ = image.shape

        (left, top, right, bottom) = (int(xmin * w), int(ymin * h),

                                     int(xmax * w), int(ymax * h))

        label = labels[classes[i]] if classes[i] < len(labels) else "Unknown"

        cv2.rectangle(image, (left, top), (right, bottom), (0, 255, 0), 2)

        cv2.putText(image, f"{label}: {scores[i]:.2f}", (left, top - 10),

                    cv2.FONT_HERSHEY_SIMPLEX, 0.6, (0, 255, 0), 2)


# Display image with detections

cv2.imshow("Detected Objects", image)

cv2.waitKey(0)

cv2.destroyAllWindows()

```

## Output :

